

Separating biological change from observation in an invasive plant time series

Working manuscript, 7 October 2026. This draft reports completed analyses and distinguishes them from planned validation. It is not ready for submission: segmentation accuracy, demographic events, the historical endpoint transformation and any evolutionary mechanism remain unvalidated. Authorship and affiliations should be supplied by the research team.

Abstract

Aerial imagery offers repeated observations of biological invasions, but changes in segmented patches need not represent changes in population growth. We reconstructed the provenance of a 2014–2021 image archive of the invasive cordgrass *Spartina alterniflora* and examined how measurement, sample selection and prediction affect its interpretation. Fifteen paired orthomosaics and masks were recovered, together with polygon annotations for 714 images containing 9,234 objects. A fixed crop convention linked 5,745 of 5,746 archived patch-date positions to mask components. All 854 records in a historical growth-model table could be traced to annual matrix transitions, preserving repeated-record multiplicity. Their inclusion was equivalent to retaining relative area increments between zero and ten from 957 transitions. Stored growth endpoints differed from area-derived relative increments in 188 records concentrated in two years. Holding the sample constant, the correlation between year and annual mean endpoint changed from -0.714 to -0.107 after recomputation. Shared mask objects, possible coverage gaps and unresolved image alignment further limited demographic interpretation. A fixed benchmark using starting geometry and purged forward-year validation did not demonstrate reliable temporal prediction. An external annual satellite product showed nonmonotonic classified cover and substantial disagreement with isolated-patch masks, identifying a further need for validation across observation scales. These findings identify the observation and response definitions that must be resolved before aerial phenotypes can support ecological mechanisms or evolutionary inference. We provide reproducible audits and a frozen, spatially separated annotation pilot for that validation.

Main text

Biological invasions alter ecosystems through changes in abundance, spatial spread and interactions with resident organisms. In coastal wetlands, *Spartina alterniflora* can transform habitat structure and ecosystem processes [1]. Repeated aerial surveys offer a way to observe these changes at the scale of individual vegetation patches. Advances in image segmentation make it feasible to extract many such observations, and spectral and phenological information has already been used to study this invasion [2]. The scientific challenge is to establish what the extracted objects measure.

A vegetation patch is an observation unit whose identity can change as plants grow, overlap or become obscured. Patch disappearance may represent death, incorporation into continuous vegetation, incomplete image coverage or a failure of detection. Likewise, a new segmented object does not necessarily identify a newly recruited genetic individual. If these processes affect inclusion in a growth dataset, the properties of retained observations may change even when the underlying population process does not. Distinguishing observation from biology is therefore necessary before interpreting temporal change.

Evolutionary explanations require an additional evidential step. Common-garden work on *S. alterniflora* has documented persistent differences between growth forms under shared conditions [3]. Studies of other invasive plants have combined common gardens, genetic information or environmental gradients to test adaptation [4,5]. Such evidence motivates evolutionary hypotheses but cannot be replaced by a correlation between year and remotely observed growth. We therefore examine the recoverable evidence chain before assigning mechanisms to the historical image series.

Reconstructing the image-to-table evidence chain

We independently rehashed all 1,001 files in the recovered archive, comprising 38,654,767,408 bytes, and found no checksum discrepancies. Fifteen paired RGB images and binary masks had identical array dimensions of 26,606 by 24,443 pixels. RGB rasters carried a shared EPSG:4326 transform, whereas masks lacked a geographic reference. Equal dimensions and header transforms establish a common array convention but do not establish the accuracy of registration.

Two annotation deliveries contained matching polygon sets after standardizing ring start, orientation and redundant closing vertices: 714 images and 9,234 objects. Polygon rasterization and 79 delivered PNG masks had a median intersection over union of 0.9965. This comparison measures consistency between annotation encodings, not the accuracy of the segmentation model. The original training photographs, segmentation checkpoint and fixed training/validation/test split have not been recovered, and the original model performance remains unverified.

A fixed region of interest, inferred from the first survey and then held constant, linked 5,745 of 5,746 archived patch-date positions to recovered mask bounding boxes. Archived x and y represented bounding-box row and column within this crop, rather than full-image centroids. An exclusive-maximum rectangular foreground count reproduced 5,744 historical areas. This rule omits the final bounding-box row and column and can include foreground belonging to nearby objects. We consequently retain historical areas as archived measurements rather than validated component areas.

The response and sample definition change the apparent temporal pattern

The annual matrix contained 2,459 candidate trajectory rows, of which 1,769 appeared in only one survey. Connecting successive observed years within rows produced 957 transitions. All had consecutive annual labels, although the selected survey dates varied within years. These

transitions are conditional on the archived tracking procedure and do not constitute a census of demographic events.

All 854 records in the modelling workbook matched annual transitions on starting year, initial area, area increment and next-survey bounding-box position, preserving the frequency of repeated keys. Four records had multiple matching candidates, so the mapping did not establish a unique biological identity for every record. The retained set was exactly reproduced by the rule $0 < (A_1 - A_0)/A_0 < 10$. This excluded 50 nonpositive changes and 53 relative increments of at least ten. The model thus described a selected set of growing, successfully tracked objects.

The stored growth_rate was inconsistent with $(A_1 - A_0)/A_0$ for all 138 records starting in 2016 and all 50 starting in 2018. The other five years agreed within numerical tolerance. No recovered transformation explained these differences. Accordingly, recomputation is a sensitivity analysis, rather than a claim to recover the intended original endpoint.

For the same 854 records, the Spearman correlation between year and annual mean endpoint was -0.714 for stored values and -0.107 for area-derived relative increments, with exact two-sided rank-permutation P values of 0.088 and 0.840, respectively (seven annual means). Including all 957 archived transitions gave a correlation of -0.607 for area-derived relative increments (exact P = 0.167). The latter mean is sensitive to extreme changes from small initial areas. These comparisons show that a mechanistic account of a temporal trend requires a resolved endpoint and an explicit sampling population; none of the comparisons independently establishes such a mechanism.

Objects, trajectories and observation gaps

Eight-connected component extraction produced 8,280 components across the 15 cropped masks, including small components. Component areas summed exactly to independently calculated mask foreground counts in each survey. Eighty component-date objects were shared by 167 archived observations, giving 87 excess rows relative to unique objects. Shared endpoints may reflect branching tracks or duplicate assignments. Their cause remains to be determined, and these observations cannot be treated as independent individuals.

An unregistered adjacent-survey overlap screen examined 5,459 source observations. Of these, 1,386 continued to an archived successor that passed the overlap criterion, and 1,123 continued to a successor that did not. Another 177 lacked a next-row observation but overlapped another mask object; 2,772 lacked both a next-row observation and a qualifying overlap. One source bounding box was unresolved. These categories are diagnostic candidates, not rates of survival, recruitment, splitting or merging.

Eighty observations without a next-row record had at least 95% exact black or white RGB values at their preceding foreground locations in the next image. Inspection of diagnostic crops confirmed that complete black coverage gaps occur. This result demonstrates an observation problem without establishing that all flagged locations are missing data. Death cannot be assigned from absence in an unvalidated mask alone. The distinction parallels the separation of

occurrence and detection in ecological observation models [6], although the current surveys do not by themselves satisfy the replicate-observation assumptions required for a fitted detection model.

Automatic feature matching did not resolve the alignment problem. Six of 14 adjacent image pairs supplied sufficient matches for a partial-affine diagnostic fit. Spatially held-out feature residuals remained substantial, and no pair passed the descriptive diagnostic gate. Features may include changing vegetation or erroneous matches; these residuals are not independent ground-control errors. No fitted transform was applied to the imagery or used to classify biological events.

Prediction with starting information and explicit holdouts

We established a fixed benchmark using log initial area and initial bounding-box coordinates, excluding next-survey locations and legacy principal components. Three exactly repeated transitions were removed, leaving 954 observations in 661 dependency groups. Groups joined annual trajectory rows sharing a dated bounding-box position. We evaluated a training-mean baseline, ridge regression and extremely randomized trees using within-fold preprocessing, forward-year prediction and buffered spatial holdouts [7].

For the primary endpoint, $\log(A_1/A_0)$, pooled forward-year R^2 values were -0.591 for the training mean, -0.941 for ridge regression and -1.116 for trees (455 predictions across 2016–2020). Performance differed by year; for example, positive R^2 in some earlier folds did not extend to 2019 and 2020. Both fitted models had lower pooled mean absolute error than the mean baseline, while their squared errors were larger. The benchmark therefore reveals sensitivity to large prediction errors and temporal transfer, rather than supporting a uniform statement about all loss functions.

Within-site buffered spatial R^2 values for log area ratio were -0.017 , 0.051 and 0.086 , respectively, over 954 predictions. For relative area increment, trees achieved a spatial R^2 of 0.278 but a forward-year R^2 of -0.601 . Initial area also enters both response definitions, so apparent size dependence need not identify a biological mechanism. These benchmarks do not evaluate climate effects or establish independent-site transfer. They identify the performance of starting-geometry baselines on archived, still unvalidated labels.

External landscape observations expose a second measurement scale

We obtained the public CCAV-10m V4 coastal vegetation product [8], including eight annual maps for 2016–2023, a sample workbook and the class-code key. All three downloaded files matched the provider's checksums. The sample workbook was byte-identical to an existing project copy, resolving its provenance. Of 13,059 public rows, 13 were exact duplicates on coordinates, year and standardized class. Within the full aerial-image footprint, 106 rows represented 99 unique coordinates, all labelled as mangroves. No local *S. alterniflora* labels were present. Training/validation membership was unavailable, so neither these records nor the source paper's reported national accuracy supplied an independent local accuracy estimate.

After nearest-neighbour reprojection to a fixed 10 m grid, the satellite product classified 97.62 ha of the historical ROI as *S. alterniflora* in 2016, 24.82 ha in 2019, 71.64 ha in 2022 and 35.22 ha in 2023 (Fig. 4). Four grid phases changed those estimates to ranges of 97.62–98.47, 24.68–25.79, 71.39–71.89 and 35.22–35.75 ha, respectively. These ranges describe discretization sensitivity, not classification uncertainty. Mangrove and unmapped areas also changed. From 2016 to 2017, 49.48 ha changed from the *S. alterniflora* class to the mangrove class, whereas 8.60 ha changed in the opposite direction. These are map-class transitions; species replacement remains an untested explanation alongside classification and observation effects.

Projecting recovered mask foreground pixel centres into the corresponding annual satellite grids yielded *S. alterniflora*-class agreement of 20.1–65.0% across 12 aerial surveys from 2016 to 2021 (Extended Data Fig. 4). Foreground totals matched independently extracted component totals in every survey. The agreement is conditional on historical mask positivity and is not an accuracy estimate. Small patches may occupy mixed 10 m pixels, annual composites differ from individual flight dates, and absolute georeferencing remains unverified. The satellite results therefore identify where landscape composition and isolated-object measurements must be reconciled, without validating either product as biological truth.

Discussion

The recovered record is valuable because it allows the observation process to be inspected. It also changes the interpretation of the historical analysis. The strongest current results concern provenance, conditioning and prediction. They do not establish a decline in total invaded area, inherited changes in growth, adaptive ecotypes or a causal effect of climate. A large calibration score cannot resolve an undefined response or an observation process that selectively removes shrinking, merging or unobservable patches.

External experimental research sharpens the competing explanations. A greenhouse study of ten Chinese and sixteen US source populations reported greater above- and belowground biomass in invasive-source plants and range-dependent relationships among latitude, flowering and biomass [9]. A three-garden study of nine Chinese seed origins reported environment-dependent germination and seedling responses [10]. The latter conclusion was assessed from the published abstract and archived dataset methods; we have not completed a full-text review or raw-data reanalysis of that experiment. Together, these studies motivate separate hypotheses about source-associated traits and environmental responses, rather than a universal direction of change in adult patch expansion. Maternal effects, ancestry, life stage and garden conditions remain relevant. Neutral-marker differentiation among Chinese and US populations [11] further motivates ancestry controls but does not establish adaptive change at this site. The two public common-garden dataset records were archived; file-access failures prevented reanalysis, and their reported results remain attributed to the original authors.

A falsifiable ecological question is whether the transition from isolated establishment to landscape occupation is governed by crowding and coalescence, neighbouring vegetation, environmental responses or heritable source differences. The satellite maps now supply candidate locations for assessing vegetation boundaries and map disagreement. Distinguishing these mechanisms requires independent species labels and event histories; assigning local

adaptation additionally requires inherited trait differences and fitness comparisons across environments. Neither a present-day common garden alone nor a space-for-time comparison would demonstrate evolution during this historical image series.

The next ecological analysis should define an explicit population at risk. A suitable representation would distinguish visible patch objects from candidate biological trajectories, permit multiple objects to merge, and record intervals in which an object cannot be observed. Growth estimates should be conditional on verified continuity, with separate treatment of disappearance and merger. Continuous vegetation should be mapped alongside isolated patches so that incorporation into a larger canopy is not mistaken for a loss of invaded area. Environmental effects should then be compared with density, neighbourhood structure and observation conditions as competing explanations.

Independent labels are required to make that representation testable. We have frozen a pilot sample of two image tiles in each of 16 geographic strata, yielding 32 locations repeated across 15 surveys. Tile selection did not consult the historical masks. Development, validation and test roles are assigned to entire spatial strata across all dates, with a buffer around held-out strata. The 480 native RGB crops and an annotation interface are available locally; completed expert labels do not yet exist. This pilot is intended to estimate annotation effort, disagreement and detection error before choosing a confirmatory sample size. It must not be presented as a completed independent validation of the historical segmenter, whose original training locations are unknown.

An evolutionary extension would require evidence that distinguishes inherited variation from plasticity and observation. Common gardens or reciprocal transplants across independent source populations, supported by clone identity or genetic data, could test whether verified field phenotypes predict persistent growth and tolerance differences. Sample size should follow variance and effect-size estimates from a pilot, with source populations—not image tiles—as the relevant independent replication. Contemporary samples cannot, without additional evidence, reconstruct genetic change between historical annual image cohorts. This limitation should remain explicit even if predictive performance improves.

The broader opportunity is to measure how invasion alters its own observable units as isolated patches expand and join continuous stands. Demonstrating that process across sites would connect AI measurement to demographic and ecological theory. The present work supplies a reproducible basis for that test and identifies the evidence still needed; it does not claim that the general mechanism has already been demonstrated.

Methods

Archive, coordinate and area audits

Source files were read without modification. SHA-256 values and byte counts were compared with the recovery manifest. Raster metadata, binary-value histograms and downsampled overlays were examined separately. The region of interest comprised rows 4,000–21,999 and columns 10,000–24,999 in the full arrays. This crop rule was inferred from the 2014 survey and

tested without refitting on subsequent dates. Bounding-box joins, external-contour counts and rectangular slice counts were retained as separate audit variables. The connected-component audit used eight-neighbour connectivity. UAV component areas and distances remain in pixels because registration and measurement accuracy are unresolved; satellite map areas are separately reported as projected classified area. The first two survey filenames specify months, not known days.

Endpoint provenance and selection

Positive annual area entries were connected within matrix rows. Relative increment was defined as $(A_1 - A_0)/A_0$; log area ratio was $\ln(A_1/A_0)$. Neither was labelled an annualized rate. Workbook-to-matrix matching preserved the multiplicity of repeated keys, and ambiguous matches were reported. Stored and recomputed endpoints were compared at tolerance 10^{-8} . Annual trend statistics used seven annual means. Exact two-sided P values enumerate all 5,040 permutations of annual ranks and assume exchangeability under the null; they do not model serial dependence. These comparisons are descriptive, not tests of heritability or environmental causation.

Transition and alignment diagnostics

Adjacent masks were linked by raw-array intersections conditional on component area ≥ 100 pixels and intersection ≥ 10 pixels. Intersection divided by smaller-component area was screened at 0.10, 0.25 and 0.50. Exact black or white next-survey RGB values at previous component pixels provided a fill diagnostic. A separate feature-alignment screen used 16 fixed 2,048-pixel windows, downsampled by two, SIFT descriptors, mutual nearest matching and a 0.65 ratio criterion. Partial-affine RANSAC used one spatial checkerboard fold; the other fold supplied diagnostic residuals. These analyses did not correct rasters or supply independently validated event labels.

Predictive benchmark

Identical transitions were deduplicated on year, both areas and both bounding-box positions. A union-find graph linked annual rows that shared a dated position, and the resulting groups were excluded from training whenever represented in a test fold. Forward folds predicted each starting year from 2016 through 2020 using earlier starting years only. Spatial folds withheld four quadrants defined at ROI column 7,500 and row 9,000, with a 512-pixel exclusion around held-out initial positions. All folds and predictions were saved.

The fixed model set comprised the training mean; median imputation, standard scaling and ridge regression with $\alpha = 10$; and median imputation followed by 160 extremely randomized trees with minimum leaf size 10. The random seed was 20261007. No hyperparameter tuning or post hoc best-model claim was made. Median imputation and scaling were fitted on each training fold. We reported pooled and fold-level R^2 , root mean squared error and mean absolute error. The only features were log initial area and the two initial coordinates. Climate, next-observation geometry, historical model outputs and precomputed principal components were excluded.

Frozen annotation pilot

The ROI was divided into 1,080 tiles of maximum side length 512 pixels, retaining partial edge tiles. Tile centres assigned membership in 16 geographic strata. Two tiles were sampled without replacement per stratum using seed 20261008; inclusion probability was $2/N_h$ for stratum h . All dates were exported at the selected locations without filtering by old-mask positivity or visibility. Sixteen locations were designated development, eight validation and eight test. A 1,024-pixel buffer around held-out strata leaves six sampled development tiles wholly eligible for training; the others must not be used for training within that buffer. The pilot does not yet provide annotated outcomes or a precision estimate. Future design-based totals must account for unequal tile inclusion probabilities, partial tiles and repeated observations of the same locations.

External satellite context and sample provenance

CCAV-10m V4 files were downloaded from Science Data Bank (10.57760/sciencedb.31077) on 7 October 2026 under CC-BY-4.0. We verified provider MD5 digests and generated SHA-256 manifests. Sample class names were stripped of surrounding whitespace. Exact duplicates were defined by longitude, latitude, year and class. WGS84 was assumed for sample coordinates because the workbook did not explicitly state a CRS. Sample membership in the product's original training and validation sets was unavailable.

National rasters were read through local windows and reprojected by nearest neighbour to EPSG:32650 at 10 m resolution. Pixel centres determined inclusion in the full-image footprint (117.404195622047–117.432930102047° E; 23.9134768474132–23.9374309174132° N) and in the historical ROI derived from the aerial header transform. Projected pixel area was 0.01 ha. Four grid origins, separated by 0 or 5 m in each direction, assessed discretization sensitivity. All class totals partitioned each footprint; adjacent-year transitions partitioned the same fixed grid. Code 255 was retained as unmapped/NoData and was not interpreted as habitat absence. No classification-error adjustment or area confidence interval was estimated.

For each recovered mask from 2016–2021, foreground pixel centres within the historical ROI were converted using the common aerial header transform and assigned the corresponding annual satellite class. Counts were checked against the independently generated component table. This conditional comparison did not sample negative masks, identify genetic individuals or estimate segmentation accuracy. National classified rasters, local clips, class keys, source digests, annual tables and analysis scripts are retained separately from the recovered originals.

Data and code availability

Code, archived tables, derived audit outputs, fixed fold membership and the pilot sampling manifest are in <https://github.com/ydchen0806/spartina-ai-eco-evolution>. The companion dataset is <https://huggingface.co/datasets/cyd0806/spartina-ai-eco-evolution-data>; a GitHub update does not imply that the same changes have been synchronized there. The 38.65 GB recovered raw archive and the 480-crop annotation pilot remain on the project server and are not bundled into the code repository. Original segmentation training images, checkpoint and split, the complete endpoint-generation implementation and the old scenario exporter remain unrecovered. No

expert pilot annotations are currently claimed. External CCAV V4 data are attributed to Yuying Li et al. (10.57760/sciencedb.31077; CC-BY-4.0); local derived rasters retain that attribution and licence. The 198.22 MB national-source download is retained on the project server, with scripts and checksums in the repository. Common-garden metadata are archived for 10.5061/dryad.2rbnz7q7 and 10.5061/dryad.j3tx95xsc (CC0); their raw files have not been obtained.

References

1. Li, B. et al. *Spartina alterniflora* invasions in the Yangtze River estuary, China: An overview of current status and ecosystem effects. *Ecological Engineering* 35, 511–520 (2009). <https://doi.org/10.1016/j.ecoleng.2008.05.013>
2. Tian, J. et al. Development of spectral-phenological features for deep learning to understand *Spartina alterniflora* invasion. *Remote Sensing of Environment* 242, 111745 (2020). <https://doi.org/10.1016/j.rse.2020.111745>
3. Gallagher, J. L., Somers, G. F., Grant, D. M. & Seliskar, D. M. Persistent Differences in Two Forms of *Spartina Alterniflora*: A Common Garden Experiment. *Ecology* 69, 1005–1008 (1988). <https://doi.org/10.2307/1941255>
4. Maron, J. L., Vilà, M., Bommarco, R., Elmendorf, S. & Beardsley, P. Rapid Evolution of an Invasive Plant. *Ecological Monographs* 74, 261–280 (2004). <https://doi.org/10.1890/03-4027>
5. Colautti, R. I. & Barrett, S. C. H. Rapid Adaptation to Climate Facilitates Range Expansion of an Invasive Plant. *Science* 342, 364–366 (2013). <https://doi.org/10.1126/science.1242121>
6. MacKenzie, D. I. et al. Estimating Site Occupancy Rates When Detection Probabilities Are Less Than One. *Ecology* 83, 2248–2255 (2002). [https://doi.org/10.1890/0012-9658\(2002\)083\[2248:ESORWD\]2.0.CO;2](https://doi.org/10.1890/0012-9658(2002)083[2248:ESORWD]2.0.CO;2)
7. Roberts, D. R. et al. Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography* 40, 913–929 (2017). <https://doi.org/10.1111/ecog.02881>
8. Li, Y. et al. CCAV-10m: an annual spatiotemporal dataset for eastern coastal China's wetland vegetation by integrating Sentinel-1/2 observations via deep learning. *Earth System Science Data* 18, 2907–2927 (2026). <https://doi.org/10.5194/essd-18-2907-2026>
9. Liu, W., Chen, X., Wang, J. & Zhang, Y. Does the effect of flowering time on biomass allocation across latitudes differ between invasive and native salt marsh grass *Spartina alterniflora*? *Ecology and Evolution* 12, e8681 (2022). <https://doi.org/10.1002/ece3.8681>
10. Liu, W., Li, L., Ye, C., Chen, X. & Zhang, Y. Temperature affects the relative importance of phenotypic plasticity and natural selection contributing to the niche breadth of invasive plants. *Oikos* 2025, e11347 (2025). <https://doi.org/10.1002/oik.11347>
11. Xia, L., Geng, Q. & An, S. Rapid Genetic Divergence of an Invasive Species, *Spartina alterniflora*, in China. *Frontiers in Genetics* 11, 284 (2020). <https://doi.org/10.3389/fgene.2020.00284>

Reference metadata and available Crossref abstracts were retrieved and checked on 7 October 2026. Full texts of references 8, 9 and 11 were additionally consulted. This does not constitute a full-text review of every reference; publisher versions and full bibliographic formatting still require final editorial checking.

Figure legends

Figure 1 | Recovered observations and provenance. Fifteen RGB/mask overlays shown at shared array indices. Colours indicate archived mask foreground. Equal array positions are an inspection convention, not proof of registration. Dense continuous vegetation is not exhaustively represented by these masks. Source:

figures/recovery_20261007/recovered_rgb_mask_contact_sheet.png.

Figure 2 | Endpoint and sample-definition sensitivity. Annual means compare stored endpoints, area-derived increments for the same model records and increments for all archived transitions. Record-level differences occur in 2016 and 2018. Pair counts show the retained set relative to the archived set. Source:

data/derived/trajectory_audit_20261007/growth_endpoint_sensitivity.png.

Figure 3 | Prediction from starting geometry. Pooled R^2 for fixed baseline models under purged forward-year and buffered within-site spatial holdouts. Primary log area ratio and secondary relative increment are shown separately. Negative R^2 denotes performance below the pooled held-out mean reference; the training-mean model is reported separately. Archived labels and identities remain unvalidated. Source:

data/derived/strict_benchmark_20261007/strict_benchmark.png.

Extended Data Figure 1 | Image alignment diagnostic. Residuals of matched image features in spatially held-out tiles before and after a diagnostic transform. Missing points denote insufficient matching support. Feature residuals do not replace independent ground-control measurements. Source: data/derived/trajectory_audit_20261007/image_alignment_diagnostic.png.

Figure 4 | Satellite classified area across fixed footprints. Annual *S. alterniflora* and mangrove class areas in the full-image footprint and historical patch ROI. Lines use the baseline 10 m UTM grid; shading spans four grid phases. The shaded range is not a classification-error confidence interval. Unmapped areas vary over time and remain separately tabulated. Source:

data/derived/ccav_20261007/ccav_classified_area.png. Derived from CCAV-10m V4, Yuying Li et al., CC-BY-4.0.

Extended Data Figure 2 | Public satellite sample audit. Geographic distribution, exact-duplicate audit and locations within the image footprint. The 106 local rows reduce to 99 coordinates, all labelled as mangroves. No independent model accuracy is estimated. Source:

data/derived/ccav_20261007/ccav_sample_audit.png. Data: CCAV-10m V4, Yuying Li et al., CC-BY-4.0.

Extended Data Figure 3 | Annual satellite classes in the image footprint. Maps share a 10 m projected grid. Grey indicates unmapped/NoData; changes between classes do not independently establish species replacement. Source:

data/derived/ccav_20261007/ccav_annual_site_maps.png. Derived from CCAV-10m V4, Yuying Li et al., CC-BY-4.0.

Extended Data Figure 4 | Agreement of isolated masks with annual satellite classes. Fractions refer to recovered ROI foreground pixel centres, assigned the same-year satellite class. Both products remain unvalidated locally; the comparison is conditional on mask positivity, includes acquisition-time and scale differences, and is not an accuracy assessment. Source: [data/derived/ccav_20261007/isolated_mask_ccav_agreement.png](#). Satellite source: CCAV-10m V4, Yuying Li et al., CC-BY-4.0.

Figure 1




Figure 2




Figure 3







Figure 4




Extended Data Figure 1



Extended Data Figure 2



Extended Data Figure 3



Extended Data Figure 4

